

# CSCE 435 Group project

---

## 0. Group number: 23

---

## 1. Group members:

---

1. Jeffrey Mitchell
2. Ren Mai
3. Brandon Cisneros
4. Christiana Vancura

## 2. Project topic (e.g., parallel sorting algorithms)

---

This project seeks to both implement and evaluate the similarities and differences between different parallel sorting algorithms with regards to problem size, problem type, and behavior regarding both strong and weak scaling. The parallel sorting algorithms chosen for the scope of this project include the following: bitonic, sample, merge, and radix sort.

### 2a. Brief project description (what algorithms will you be comparing and on what architectures)

#### **Bitonic Sort:**

Bitonic sort assumes an input size that is a power of 2. This algorithm uses many different processors as a network of comparators in order to sort a sequence. This algorithm must first ensure that data is in a bitonic sequence before sorting. A sequence of numbers is bitonic if it consists of a first substring of strictly non-decreasing numbers, followed by a substring of strictly non-increasing numbers. Once a bitonic sequence is generated, the algorithm can then merge the sequence into a fully sorted list. This can be repeated at scale to sort large sequences. To implement this sorting algorithm, we will be using the MPI architecture.

#### **Sample Sort:**

Sample sort does not assume any input size, but will split the input into  $p$  (the number of processors) buckets. Thus it will be easier to use inputs of size powers of 2. The algorithm splits the initial input into  $p$  sections, sorts them using quicksort, and samples  $s$  elements from each. These samples are then sorted (again using quicksort), and  $p-1$  elements are chosen as splitters. Each processor splits its input bucket based on the splitters and sends the resultant buckets to the corresponding processor (bucket 0 to 0, 1 to 1, and so on). The total input is thus split into  $p$  buckets, which are sorted using bucket sort by each processor. These are then combined to return a sorted list. We will be using MPI architecture to implement sample sort.

### **Merge Sort:**

Merge sort works by repeatedly merging larger and larger sorted subsets of an array. It is not confined to particular input array sizes, and as a comparison-based sorting algorithm, it can operate on any comparable data type. Merge sort is a divide and conquer algorithm. In its base case, a subarray of length 0 or 1 is already sorted, and at each step, merge sort merges two sorted arrays by stepping through them and popping the smallest remaining element from either. In a parallel environment, each of the  $p$  nodes will sort itself normally, then will cooperate to continue the merges across process boundaries, until the array is fully sorted. For the first layer, odd and even processes will merge with each other, with the odd processes taking the higher halves of each pair's array elements. The processes will then sort themselves into sets of four and use three 2-way merges to generate consistency among the four processors. The number of processes per group will continue doubling until the full composite array is sorted. Our implementation will use the MPI architecture.

### **Radix Sort:**

Radix sort assumes that the input elements of the problem statement are  $k$  digit numbers. The algorithm sorts the elements through means of buckets; by sorting elements by their digits, Radix Sort is able to linearly sort elements as it sorts from least to most significant digit. For the sake of this sorting algorithm, we will be using the MPI architecture.

## **2b. Pseudocode for each parallel algorithm**

- For MPI programs, include MPI calls you will use to coordinate between processes

### **Bitonic Sort:**

**MPI calls used to coordinate between processes:**

- `MPI_Init(...)`
- `MPI_Comm_rank(...)`
- `MPI_Comm_size(...)`
- `MPI_Send(...)`
- `MPI_Recv(...)`
- `MPI_Comm_split(...)`
- `MPI_Gather(...)`
- `MPI_Scatter(...)`
- `MPI_Finalize(...)`
- `MPI_Barrier(...)`

## Pseudocode

```

if worker process:
    receive subarray

    # Make own subarray bitonic sequence
    for each 2-element pair in subarray:
        if odd count pair:
            sort ascending
        if even count pair:
            sort descending
    for each 4-element pair in subarray:
        if odd count pair:
            sort ascending
        if even count pair:
            sort descending
    etc. until subarray is bitonic

    # Multi-step bitonic merge
    for each compare-swap level needed (1 - log2(# of processes)):
        find partner process
        exchange data with partner (MPI_sendrecv)
        if "inferior" partner:
            keep smaller half of data
        else if "superior" partner:
            keep larger half of data

    Send sorted subarray to master

else if master process:
    send subarrays to worker processes (MPI_Scatter)
    receive sorted subarrays/gather them (MPI_Gather)

```



```
merge sorted subarrays
return sorted list
```

## Sample Sort:

### MPI calls used to coordinate between processes:

- MPI\_Init(...)
- MPI\_Comm\_rank(...)
- MPI\_Comm\_size(...)
- MPI\_Send(...)
- MPI\_Recv(...)
- MPI\_Comm\_split(...)
- MPI\_Gather(...)
- MPI\_Scatter(...)
- MPI\_Allgather(...)
- MPI\_Finalize(...)
- MPI\_Barrier(...)

### Pseudocode

```
split initial array into subarrays
for each process
    sort each subarray using quicksort
    send process 0 `s` elements (MPI_Send)
barrier (MPI_Barrier)
process 0 receives `s` elements (MPI_Gather)
process 0 sorts elements with quicksort
process 0 chooses `p-1` splitters
process 0 sends splitters to all processes (MPI_Scatter)
for each process
    receive splitters (MPI_Receive)
    from each subarray split on splitters into `p` buckets
    send bucket 0 to process 0, bucket 1 to process 1, ..., bucket `p-1`
    barrier (MPI_Barrier)
    processes receive buckets (MPI_Allgather)
    bucket sort received buckets
    send received buckets to process 0 (MPI_Send)
process 0 receives and concatenates buckets (MPI_Gather)
process 0 returns sorted bucket 0, bucket 1, ..., bucket `p-1`
```



## Merge Sort:

### MPI calls used:

- `MPI_Init(...)`
- `MPI_Comm_size(...)`
- `MPI_Comm_rank(...)`
- `MPI_Send(...)`
- `MPI_Recv(...)`
- `MPI_Finalize(...)`

### Pseudocode

// Sort a fully-local array with standard merge sort:

```
local_merge_sort(array, size):
    if (size < 2):
        return
    local_merge_sort(array, size/2)
    local_merge_sort(&array[size/2], size/2)

    tmp_array_i = new array of length size/2
    tmp_array_j = new array of length size/2
    for i from 0 to size/2 - 1, inclusive:
        tmp_array_i[i] = array[i]
        tmp_array_j[i] = array[size/2 + i]

    i = 0
    j = 0
    while (i < size/2 or j < size/2):
        if (i == size/2):
            array[i+j] = tmp_array_j[j]
            j++
            continue
        if (j == size/2):
            array[i+j] = tmp_array_i[i]
            i++
            continue
        if (tmp_array_i[i] <= tmp_array_j[j]):
            array[i+j] = tmp_array_i[i]
            i++
        else:
            array[i+j] = tmp_array_j[j]
            j++

// Take only the bottom half of the provided values and merge them:
merge_bottom_half(array_i, array_j, array_out, size):
```



```

i = 0
j = 0
while (i + j < size):
    if (array_i[i] <= array_j[j]):
        array_out[i+j] = array_i[i]
        i++
    else:
        array_out[i+j] = array_j[j]
        j++

// Take only the top half of the provided values and merge them:
merge_top_half(array_i, array_j, array_out, size):
    i = size-1
    j = size-1
    out_place = size-1
    while (out_place >= 0):
        if (array_i[i] <= array_j[j]):
            array_out[out_place] = array_i[i]
            i--
        else:
            array_out[out_place] = array_j[j]
            j--
        out_place--

// Merge with one specific neighbor. Relative ranks determine which one
// or lower array elements.
merge_2_way(neighbor_id):
    Send local_subarray to neighbor_id via MPI_Send()
    Receive remote_subarray from neighbor_id via MPI_Recv()
    new_subarray = allocate new array of length (n/p)
    if (neighbor_id > local_rank):
        // Neighbor has a higher rank, meaning this process will get the
        merge_bottom_half(local_subarray, remote_subarray, new_subarray,
    else:
        // Neighbor has a lower rank, meaning this process will get the
        merge_top_half(local_subarray, remote_subarray, new_subarray, n,
    Deallocate local_subarray
    Set local_subarray = new_subarray

// Use a combination of 2-way merges to combine two sorted chunk_size/2
// into one sorted chunk_size sized chunk. Each offset=x iteration isolat
// process at the top and bottom of the chunk as definitely being sorted
// new chunk. At the last iteration, only the middle two processes merge
// other, because they're the only two not confirmed to be sorted yet.
merge_n_way(chunk_size):
    rank_within_chunk = local_rank % chunk_size

    offset = chunk_size / 2
    min_offset = chunk_size/2 - rank_within_chunk

```

```

    if (rank_within_chunk >= chunk_size / 2):
        min_offset = -min_offset + 1
    while (offset >= min_offset):

        // Choose offset direction based on whether this process is in 1
        // half of the chunk:
        if (rank_within_chunk < chunk_size/2):
            merge_2_way(local_rank + offset)
        else:
            merge_2_way(local_rank - offset)
        offset--

MPI_Init()

p = MPI_Comm_size()
local_rank = MPI_Comm_rank()

local_subarray = New array with (n/p) elements generated according to th

local_merge_sort(local_array, n/p)
k = 1
while k < n:
    k *= 2
    // Example: k=4.
    // Odd and even processes have already merged, meaning there are p/2
    // subarrays spanning 2 processes each, and this next step will cons
    // sorted distributed subarrays spanning 4 processes each.
    merge_n_way(k)

// Double-check sorting:
for i from 1 to (n/p - 1):
    assert local_array[i-1] <= local_array[i]

// Send top value to next process and bottom value to previous process:
if (local_rank < p-1):
    Send local_array[n/p - 1] to (local_rank+1) via MPI_Send()
if (local_rank > 0):
    Send local_array[0] to (local_rank-1) via MPI_Send()
    Receive the value next_lower from (local_rank-1) via MPI_Recv()
    assert local_array[0] >= next_lower
if (local_rank < p-1):
    Receive the value next_higher from (local_rank+1) via MPI_Recv()
    assert local_array[n/p - 1] <= next_higher

MPI_Finalize()

```


## Radix Sort:

## MPI calls used to coordinate between processes:

- `MPI_Init(...)`
- `MPI_Comm_rank(...)`
- `MPI_Comm_size(...)`
- `MPI_Send(...)`
- `MPI_Recv(...)`
- `MPI_Comm_split(...)`
- `MPI_Gather(...)`
- `MPI_Finalize(...)`
- `MPI_Barrier(...)`

## Pseudocode

**Note:** Implementation can take many forms with regards to digits. For example, we can do base 10, base 2, and etc. For the sake of this pseudocode, we abstract this away by simply calling the extracted variable `digit`.

**# Note:** we can either generate the problem using centralized master, and 

**if** Master:

generate the specified problem `type`(sorted, sorted with `1%` swap, etc)  
send offsets to worker processes `#MPI_Send(...)`

**if** worker process:

receive offset for processing `#MPI_Send(...)`  
**#create** local bucket storage  
`histogramBucket[numProcs][numBuckets]`

**# compute** and build histogram for offset (sort local offset) by dete

**for** each item `i` in the array offset:  
  **#determine** bucket to place `a[i]`  
  `bucket = compute_bucket(a[i])`  
  **#increment** the bucket count  
  `histogramBucket[numProc][bucket]++`

`MPI_Barrier()` **# to** wait for all worker processors to finish computin

**# master** processor determines the position for each by doing prefix sum  
**# this** is the sequential portion

**if** master proc:

`base = 0`  
  **for** bucket in `numBuckets`:



```

for numProc in numProcs:
    # add base to the histogramBucketLocation and compute partial sum
    histogramBucket[numProc][bucket] += base

    #update base, we are computing prefix
    base = histogramBucket[numProc][bucket]

MPI_Barrier(...) # wait for this indexer algorithm to be determined
MPI_Scatter(...) # send the indices information to the workers so they can
if worker process:
    for each item i in array offset:
        bucket = compute_bucket(a[i])
        #place the item a[i] in the correct location
        # send the item to the proc that takes care of that offset or handle
        if histogramBucket[numProc][bucket] index is within the worker's range:
            outputResult[histogramBucket[numProc][bucket]++] = a[i]
        else:
            MPI_send(...) to the correct proc with the index
MPI_Recv(...) items and their indices and place into correct offset

```

## 2c. Evaluation plan - what and how will you measure and compare

### Input sizes, Input types

With regards to input sizes and input types for the project, the sorting algorithms implemented will test the following specifications:

#### Input sizes for the array being sorted:

The input size of the array being sorted will take on sizes such that size corresponds to powers of two. This is done to ensure that data partitioning for each process will result in relatively similar problem subset sizes. In particular, the problem sizes being tested involve the following:  $2^{16}$   $2^{18}$   $2^{20}$   $2^{22}$   $2^{24}$   $2^{26}$   $2^{28}$

#### Input types:

Input types utilized in the project will take the following configurations:

- Sorted
  - The array passed in as input will already be in sorted form.
- Random

- The array passed in as input will include randomly generated elements such that the array is very likely unsorted.
- Reverse sorted
  - The array passed in as input is sorted, but uses the opposite order.
  - *(e.g. If we are sorting descending, the reverse sorted array will be ascending and vice versa)*
- 1% Perturbed
  - The array is mostly sorted, with 1% of array items being swapped or placed in unsorted positions.

### **Number of Processes used:**

For the sake of this project, the number of processes being tested will also be done using powers of two. More specifically, the process counts utilized in runs will include the following process counts: 2, 4, 8, 16, 32, 64, 128, 256, 512, and 1024.

### **Strong scaling**

The strong scaling potential of an algorithm can be analyzed by keeping the problem size fixed while increasing the number of processors/nodes. For each algorithm, this will be measured by recording the time it takes to work through an input of a constant size when utilizing varying amounts of processors. We will increase processor count progressively, testing powers of two for their performance (2, 4, 8, 16 processors, etc.). The actual problem size will be decided based on benchmarking a small processor count in order to ensure that jobs can complete on a reasonable timescale (hours at most). The corresponding decrease (or increase) in execution time will allow us to measure the relative strong scaling of each algorithm.

### **Weak scaling (increase problem size, increase number of processors)**

An algorithm's weak scaling is dependent on how well additional processes allow it to solve larger problems in a similar amount of time. For each of the above algorithms, we will evaluate weak scaling by measuring runtime while the number of processes and length are kept proportional: We will test a problem size  $n = 2^{20}$  with  $p = 4$  processes,  $n = 2^{22}$  with  $p = 16$ ,  $n = 2^{24}$  with  $p = 64$ ,  $n = 2^{26}$  with  $p = 256$ , and  $n = 2^{28}$  with  $p = 1024$ . When graphed, effective weak scaling will maintain a mild or even slope, and ineffective weak scaling will show steeply increasing runtime as the problem size and processor count increase.

## **2d. How the team will communicate:**

- For the sake of this project, the team has decided to go forward with using Slack as the main form of communication, with periodic in-person meetings for discussion and implementation of the algorithms at hand.